

Homework 3 (Python version)

ME570 - Prof. Tron

2023-10-14

Problem 1: Drawing and collision checking for spheres

In this problem you will write functions that are similar to those in Problem 1 of Homework 1, but applied to 2-D spheres (i.e., circles).

Question provided 1.1. A class for spheres.

File name: `me570_geometry.py`

Class name: `Sphere`

Description: Class for plotting and computing distances to spheres (circles, in 2-D).

Method name: `__init__`

Description: Save the parameters describing the sphere as internal attributes.

Input arguments

- `center` (dim. $[2 \times 1]$): the 2-D coordinate of the center of the sphere;
- `radius`: a scalar whose absolute value is equal to the geometric radius of the sphere, and the sign indicates whether the obstacle should be interpreted as filled-in (`radius` > 0) or hollow (`radius` < 0);
- `distance_influence`: a scalar containing the *influence* distance of the sphere (the exact meaning of this will become clear after we talk about potential-based planning).

Figure 1 demonstrates the meaning of these quantities.

Question provided 1.2. Implement a function that draws a sphere.

Class name: `Sphere`

File name: `me570_geometry.py`

Method name: `plot`

Description: This function draws the sphere (i.e., a circle) of the given radius, and the specified color, and then draws another circle in gray at the influence radius.

Input arguments

- `color` (dim. $[1 \times 1]$, type `string`): a color specification string (e.g., `'b'`, `'r'`, etc.)

Question code 1.1. Implement the following function.

Class name: `Sphere`

File name: `me570_geometry.py`

Method name: `distance`

Description: Computes the signed distance between points and the i -th sphere, while taking into account whether the sphere is hollow or filled in.

Input arguments

- `points` (dim. $[2 \times \text{nb_points}]$): an array of 2-D points.

Output arguments

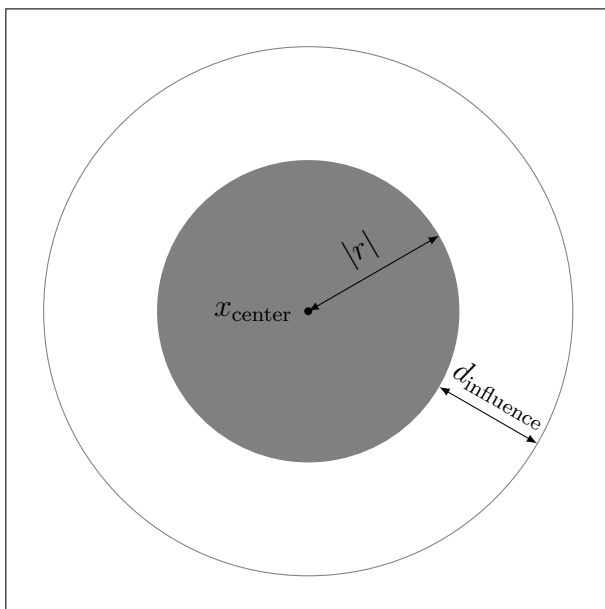
- `d_points_sphere` (dim. $[1 \times \text{nb_points}]$): the distance between each point and the surface of the sphere. The distance is negative if `points` is inside the obstacle (i.e., inside the sphere for filled-in obstacles, and outside for hollow obstacles), and positive otherwise.

Requirements: Remember that the radius of the sphere is negative for hollow spheres.


Question code 1.2. Implement a function to compute the gradient of the sphere.

Class name: `Sphere`

File name: `me570_geometry.py`



(a) Filled sphere (`radius` > 0)



(b) Hollow sphere (`radius` < 0)

Figure 1: Examples of spheres.

Method name: `distance_grad`

Description: Computes the gradient of the signed distance between points and the sphere, consistently with the definition of `Sphere.distance(.)`.

Input arguments

- `points` (dim. $[2 \times \text{nb_points}]$): an array of 2-D points.

Output arguments

- `grad_d_points_sphere` (dim. $[2 \times \text{nb_points}]$): the gradient of the distance. If a point corresponds to the center of the sphere, return the zero vector.

Question optional 1.1. Implement a function that shows collision checks with a sphere.

File name: `me570_hw3.py`

Method name: `sphere_test_collision`

Description: Generates one figure with a sphere (with arbitrary parameters) and `nb_points=100` random points that are colored according to the sign of their distance from the sphere (red for negative, green for positive). Generates a second figure in the same way (and the same set of points) but flipping the sign of the radius `r` of the sphere.

Red and green points represent locations that are, respectively, in collision and not in collision with the spherical obstacle.

The world

The goal of this question is to prepare and visualize the *sphere world* workspace that will be used in the other problems. The workspace is described by the data stored in the file `sphereworld.mat` provided with this homework.

File name: `me570_potential.py`

Class name: `SphereWorld`

Description: Class for loading and plotting a 2-D sphereworld.

Method name: `__init__`

Description: Load the sphere world from the provided file `sphereworld.mat`, and sets the following attributes:

- `world`: a `nb_spheres` list of `Sphere` objects defining all the spherical obstacles in the sphere world.
- `x_start`, a $[2 \times \text{nb_start}]$ array of initial starting locations (one for each column).
- `x_goal`, a $[2 \times \text{nb_goal}]$ vector containing the coordinates of different goal locations (one for each column).

Question provided 1.3. Implement a function to visualize a sphere world

Class name: `SphereWorld`

File name: `me570_potential.py`

Method name: `plot`

Description: Uses `Sphere.plot(_)` to draw the spherical obstacles together with a `*` marker at the goal location.

Use this function to visualize the contents of the file.

Question provided 1.4. The following is a utility function.

Class name: `me570_geometry.py`

File name: `Grid`

Method name: `plot_threshold`

Description: The function evaluates the function `f_handle` on points placed on a regular grid.

Input arguments

- `f_handle` (type `function`): an handle to a function that takes a vector `x` of dimension $[2 \times 1]$ as a single argument, and returns a scalar or a vector. An example of this function could be the function `norm(_)`.
- `threshold` **optional**, default `10`: if a vector `fEval`, obtained by evaluating `fHandle`, has `norm(fEval)>threshold`, then it is replaced by `fEval/norm(fEval)*threshold`. If omitted, `threshold = 10`.

Requirements: The function is build upon the function `Grid.eval(_)` from Homework 2.

Note that you can look at the source code of this function for a way to handle optional arguments.

Note on the use with functions that take extra parameters You can use `Grid.plot_threshold(_)` also with functions that require additional parameters (e.g., `Potential.attractive(_)`). This can be done by creating a *local function* that calls the desired function with the additional parameters are fixed. The best way to understand this is with an example:

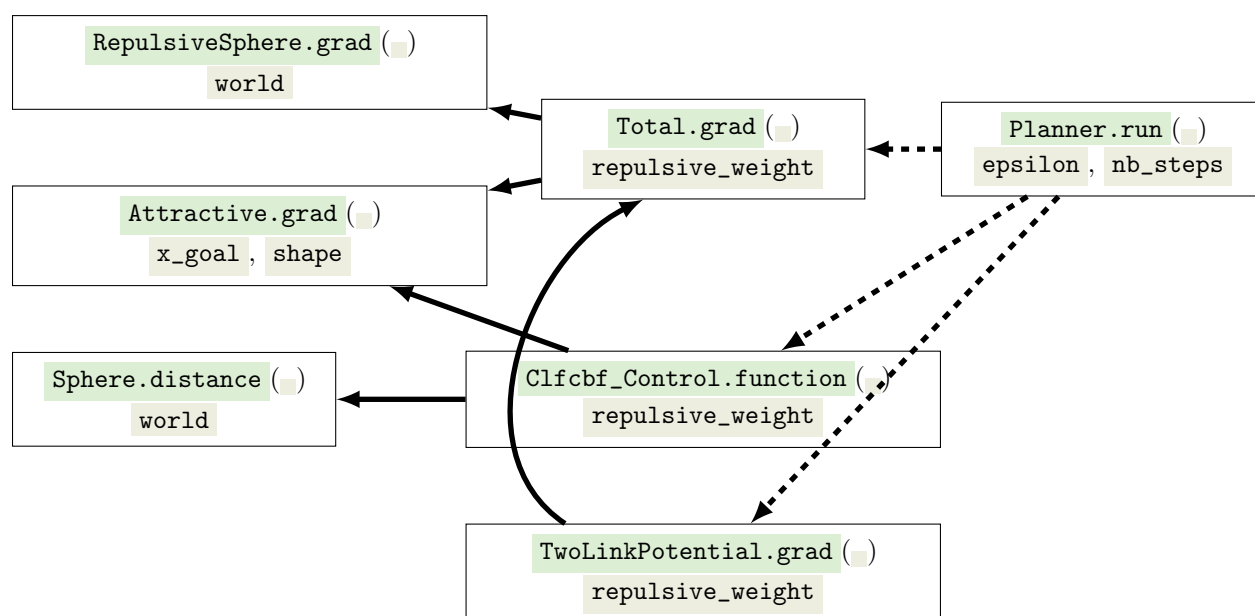
```
xx_ticks=np.linspace(-11,11,51)
grid=Grid(xx_ticks,xx_ticks)
center=np.ones((2,1))
sphere=Sphere(center,-2,3)
def f_handle(point):
    return sphere.distance(point)
grid.plot_threshold(f_handle,10)
```

Note: You should avoid defining functions that *capture* variables (e.g., `center` in the example above) that change in a loop (they might lead to some unexpected behaviors). *You need to type the example, copying from this PDF document will not work.*

The first line of the example creates an object `sphere` from the class `Sphere`. The second line creates a lambda function with a single argument `point`, and stores it in the variable `f_handle`; since `sphere` does not appear in the argument list of the lambda function, it is “captured” from the environment (i.e., it is “remembered” inside the lambda function). Now `f_handle` can be used as a function with a single input argument (e.g., try `f_handle(np.zeros((2,1)))`). In the third line, this function is passed to `grid_plot_threshold` for plotting. Note that this will still work even though `sphere` is not visible in scope inside of `grid_plot_threshold` (as `sphere` is “remembered” inside the lambda function, this is not necessary).

Organization of the reminder of the assignment

Below you can find below a diagram of the most important functions in the reminder of the homework. Solid arrows denote functions that directly call each other; dashed arrows denote functions that are called by passing them as arguments. Each box contains also the main parameters used by the function.



The code is written in such a way that many parts are reused. We will have a function that runs the potential planner, i.e., that applies Euler’s method, for three different settings. First, the traditional potential methods (Problem 2), where the planner uses the total potential which is computed as the sum of repulsive and attractive potentials. Second, the CLF-CBF method (Problem 3), where the planner uses a vector field computed, from the attractive gradient and the sphere distance function. Third, an inverse kinematics solver for a two-link manipulator (Problem 4) where the planner uses the total potential attached to the end effector of the manipulator.

Problem 2: The potential-based planner

In this problem, you will implement and test the path planning strategy based on attractive-repulsive potential. Different functions will use similar dictionaries, which are explained here for ease of reference:

- A dictionary `potential` with elements:
 - `potential['x_goal']`: a $[2 \times 1]$ vector with the x, y coordinates of the goal location.
 - `potential['repulsive_weight']`: a scalar containing the weight of the attractive term with respect to the repulsive term.
 - `potential['shape']`: a string that specifies the shape of attractive potential. It can be equal to `'conic'` or `'quadratic'`.

How these structures are actually used is explained in the questions below.

The attractive and repulsive potential. In the following, we use the shorthand notation $d_i(x)$ to denote the distance between the point x and the sphere (defined as seen in class).

The attractive potential we use for this assignment is given by:

$$U_{\text{attr}}(x) = d^p(x, x_{\text{goal}}) = \|x - x_{\text{goal}}\|^p, \quad (1)$$

where p is a parameter to distinguish between conic and quadratic potentials, and its corresponding gradient is:

$$\nabla U_{\text{attr}}(x) = p d^{p-1}(x, x_{\text{goal}}) \frac{x - x_{\text{goal}}}{\|x - x_{\text{goal}}\|} = p d^{p-2}(x, x_{\text{goal}})(x - x_{\text{goal}}), \quad (2)$$

The repulsive potential we use for this assignment is given by

$$U_{\text{rep},i}(x) = \begin{cases} \frac{1}{2} \left(\frac{1}{d_i(x)} - \frac{1}{d_{\text{influence}}} \right)^2 & \text{if } 0 < d_i(x) < d_{\text{influence}}, \\ 0 & \text{if } d_i(x) > d_{\text{influence}}, \\ \text{undefined} & \text{otherwise,} \end{cases} \quad (3)$$

and the corresponding gradient is

$$\nabla U_{\text{rep},i}(x) = \begin{cases} - \left(\frac{1}{d_i(x)} - \frac{1}{d_{\text{influence}}} \right) \frac{1}{d_i(x)^2} \nabla d_i(x) & \text{if } 0 < d_i(x) < d_{\text{influence}}, \\ 0 & \text{if } d_i(x) > d_{\text{influence}}, \\ \text{undefined} & \text{otherwise.} \end{cases} \quad (4)$$

Question provided 2.1.

Class name: `RepulsiveSphere`

File name: `me570_potential.py`

Description: Repulsive potential for a sphere

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `sphere` (type `Sphere`): an object describing a sphere as described above.

Question **provided** 2.2. This function evaluates the repulsive potential.

Class name: `RepulsiveSphere`

File name: `me570_potential.py`

Method name: `eval`

Description: Evaluate the repulsive potential from `sphere` at the location $x = \mathbf{x_eval}$. The function returns the repulsive potential as given by (3).

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the total repulsive potential.

Output arguments

- `u_rep` (dim. $[1 \times 1]$): the value of U_{rep} evaluated at $x = \mathbf{xEval}$.

Use the value *Not-a-Number* (`math.nan` in Python) for the case where the potential is undefined.

Question **code** 2.1. For this question you need to implement the gradient for the repulsive potential for spheres.

Class name: `RepulsiveSphere`

File name: `me570_potential.py`

Method name: `grad`

Description: Compute the gradient of U_{rep} for a single sphere, as given by (4).

Input arguments

- `x_eval` (dim. $[2 \times 1]$): see input arguments for `RepulsiveSphere.eval(_)`.

Output arguments

- `grad_u_rep` (dim. $[2 \times 1]$): the gradient of U_{rep} evaluated at $x = \mathbf{xEval}$.

Requirements: This function should use `self.sphere.distance_grad(_)`

Use the value *Not-a-Number* (`Nan` in Matlab) for the case where the potential is undefined.

Question **report** 2.1 (0.5 points). Use the provided function `Grid.plot_threshold(_)` to visualize `RepulsiveSphere.grad(_)` for the first two spheres in the sphere world, and overlap the plot of the sphere in each plot; use a different figure for each sphere. Include the figures in your report.

Question report 2.2 (0.5 points). In the figures generated with the previous question, should the arrows point toward or away from the obstacle? Why? Why are the arrow only visible around the contour of the obstacle?

Question provided 2.3.

Class name: `Attractive`

File name: `me570_potential.py`

Description: Repulsive potential for a sphere

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `potential` (dim. $[1 \times 1]$, type `dict`): the dictionary with fields `'x_goal'`, `'shape'`.

Question provided 2.4. Implement the attractive potential

Class name: `Attractive`

File name: `me570_potential.py`

Method name: `eval`

Description: Evaluate the attractive potential U_{attr} at a point `xEval` with respect to a goal location `potential.xGoal` given by the formula: If `potential.shape` is equal to `'conic'`, use $p = 1$. If `potential.shape` is equal to `'quadratic'`, use $p = 2$.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the potential.

Output arguments

- `u_attr` (dim. $[1 \times 1]$): the value of U_{attr} evaluated at $x = \text{xEval}$.

Question code 2.2 (2 points). Implement a function to compute the gradient of the attractive potential.

Class name: `Attractive`

File name: `me570_potential.py`

Method name: `grad`

Description: Evaluate the gradient of the attractive potential ∇U_{attr} at a point `xEval`. The gradient is given by the formula. If `potential['shape']` is equal to `'conic'`, use $p = 1$; if it is equal to `'quadratic'`, use $p = 2$.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the potential.

Output arguments

- `grad_u_attr` (dim. $[2 \times 1]$): the value of ∇U_{attr} evaluated at $x = \text{xEval}$.

Question provided 2.5. The following class combines the attractive and repulsive potentials into the total potential.

Class name: `Total`

File name: `me570_potential.py`

Description: Combines attractive and repulsive potentials

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `world` (type `SphereWorld`): an instance of the class `SphereWorld` (see the corresponding description for details).
- `potential` (type `dict`): the dictionary with fields `'x_goal'`, `'shape'` and `'repulsive_weight'` previously described. This function uses the field `'repulsive_weight'`, and then the structure is passed to an instance of `Attractive`.

Question code 2.3. The following two functions should combine the attractive and repulsive potentials into the total potential.

Class name: `Total`

File name: `me570_potential.py`

Method name: `eval`

Description: Compute the function $U = U_{\text{attr}} + \alpha \sum_i U_{\text{rep},i}$, where α is given by the variable `potential['repulsive_weight']`.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the point where to evaluate the total potential.

Output arguments

- `u_eval`: The value of the potential U evaluated at `x_eval`.

Class name: `Total`

File name: `me570_potential.py`

Method name: `grad`

Description: Compute the gradient of the total potential, $\nabla U = \nabla U_{\text{attr}} + \alpha \sum_i \nabla U_{\text{rep},i}$, where α is given by the variable `potential['repulsive_weight']`

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the point where to evaluate the gradient of the total potential.

Output arguments

- `grad_u_eval` (dim. $[2 \times 1]$): The gradient of the potential, ∇U , evaluated at `x_eval`.

Question optional 2.1. Use the function `grid_plotThreshold` (■) to visualize the attractive and repulsive potentials (first separately for a goal and obstacle, and then combined in the total potential), and their gradients, for different situations (try both conic and quadratic potentials, and both filled-in and hollow spherical obstacles). For the total potential, overlap the output of `grid_plotThreshold` (■) on top of the output of `SphereWorld.plot` (■).

Question code 2.4. Implement a Class for your main potential planner.

File name: `me570_potential.py`

Class name: `Planner`

Description: A class implementing a generic potential planner and plot the results.

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `function` (type `function`): handle to the function for computing the value of the potential function;
- `control` (type `function`): handle to the function for computing the direction in which the planner should move (for gradient-based methods, this is the negative gradient of the potential function);
- `epsilon`: value for the step size;
- `nb_steps`: total number of steps.

Class name: `Planner`

File name: `me570_potential.py`

Method name: `run`

Description: This function uses a given function (given by `control`) to implement a generic potential-based planner with step size `epsilon`, and evaluates the cost along the returned path. The planner must stop when either the number of steps given by `nb_steps` is reached, or when the norm of the vector given by `control` is less than $5 \cdot 10^{-3}$ (equivalently, `5e-3`).

Input arguments

- `x_start` (dim. $[2 \times 1]$): starting location.

Output arguments

- `x_path` (dim. $[2 \times \text{nb_steps}]$): sequence of locations generated by the planner, such that `x[:,1]=x_start` and `x[:,k+1]=x[:,k]+epsilon*control_current`, where `control_current` is obtained by evaluating `control` using `x[:,k]`, `world` and `potential` as arguments. If the planner stops for some `k<nb_steps`, the remaining entries should be filled with `NaN` (the special value *Not-a-Number*).
- `u_path` (dim. $[1 \times \text{nb_steps}]$): The array of values of the potential U evaluated at each of the points on the path (that is, `u_path[k]` is the value of U evaluated at `x_path[k]`). Again, use `NaN`'s to fill the array if the planner stops before `nb_steps`.

Note: Questions report 2.3–report 2.7 should be considered together.

Question report 2.3. Implement a function that runs the planner on the provided data, and visualizes the results to explore and illustrate the effect of the different parameters in the planner. This question will generate many figures. Include all figures in your report, trying to have all of them organized in one or two pages, for ease of comparison.

File name: `me570_hw3.py`

Method name: `planner_run_plot_test`

Description: The function should call prepare a *Planner* object and call its method `run` () for every combination of start and goal location in `world`, and for different interesting combinations of the parameters. See below for the specific steps.

This function performs the following steps for every combination of `epsilon`, `shape`, `nb_steps`:

- 1) Load the problem data via an object `world` of class `SphereWorld`.
- 2) For every goal location in `world.x_goal`:
 - (a) Create a new figure, with two subplots. In the subplot on the left, use the function `Sphereworld.plot` () to plot the world.
 - (b) For every start location in `world.x_start`:

- i. Prepare the `potential` dictionary with the following fields set for the current test combination.
 - `'x_goal'`
 - `'repulsive_weight'`
 - `'shape'`
 - ii. Create an object of class `Total` using the `potential` dictionary above.
 - iii. Prepare the object of class `Planner` with the following fields:
 - attribute `function`, set to `Total.eval`
 - `control`, set to the negative of `Total.grad`.
 - `epsilon`.
 - `nb_steps`.
 - iv. Calls the method `run()`.
 - v. Plot the resulting trajectory superimposed to the world in the first subplot.
 - vi. Plot the corresponding value of the potential `u_path` (using the same color and using the `semilogy` command) in a second subplot.
- (c) For the current goal, you should have one figure with two subplots with the data for five paths.
- (d) If some of the trajectories converge, add another figure zoom in closely around the final equilibrium (to see the behavior of the planner near it).
- 3) For the current combination of parameters, you should have two figures, each one with two subplots.

Suggestions for tuning parameters. Start with `'repulsive_weight'` in the interval $0.01 - 0.1$, ϵ in the range $1e-3 - 1e-2$, for `nb_steps`, use `100`, and explore from there. Typically, adjustments in `repulsive_weight` require subsequent adjustments in `epsilon`; if the trajectories seem to behave well but stop short of the goal, increase `nb_steps`. In general, you should be able to tune the parameters such that almost all the trajectories reach the respective goal (or its vicinity). Additional hints are available for this question.

Question report 2.4. Use the function `grid_plotThreshold()` to visualize the total potential U , and its gradient ∇U in two separate figures for each one of the combinations `'repulsive_weight'` and `'shape'` included in the previous question (report 2.3) (for this question, it is sufficient to consider only one of the two goals). Include the images in your report.

Question report 2.5. Comment on the effects of varying each one of the parameters `'repulsive_weight'` and `'epsilon'`. Explain why the planner behaves in the way you see, explicitly explaining what happens for the four cases where the two parameters are, respectively, small/small, small/large, large/small, and large/large. Additionally, comment on whether the results you see in practice are consistent with what discussed in class.

Question report 2.6 (0.5 points). Comment on the relation between the value of the potential U toward the end of the iterations, versus the fact that the planner correctly

succeeded or failed. Explain different reasons why the planner might fail, and why changing the various parameters can improve the situation.

Question report 2.7. What is the difference between the two goals included in the provided dataset? In relation to this, what is the effect of the parameter `'shape'`?

Problem 3: CLF-CBF formulation

This problem is similar to Problem 2, except that you will use a CLF-CBF formulation instead of a traditional gradient-based formulation. We will use U_{attr} in (1) as the CLF, and $d_i(x)$ (implemented in `Sphere.distance` ()) as the CBF for each obstacle. Note that, since we control the position of our (point) robot directly, the dynamics of our system is input-affine with the form $\dot{x} = f(x) + g(x)u = u$, i.e., $f(x) = 0$ and $g(x) = I$, the identity matrix.

Preparation. First, we need to write the Quadratic Program (QP) that uses the minimum-effort CLF-CBF formulation (the supervisor version) to return a control $u^*(x)$ for any given x (see also the class notes). In particular, the QP will have the form

$$u^*(x) = \underset{u}{\operatorname{argmin}} \|u - \clubsuit\|^2 \quad (5a)$$

$$\text{subject to } \spadesuit \leq 0, \quad (\text{CBF constraint, } i\text{-th obstacle}) \quad (5b)$$

Question report 3.1. In the report, write the expressions for \clubsuit (involving U_{attr}) and \spadesuit (involving $d_i(x)$). For this and the questions below, denote the constant appearing in the constraint (5b) as c_h . Please pay attention to the directions of the inequalities.

Question provided 3.1. As detailed in the Week 6 in-class activity, the provided `me570_qp_qp_supervisor` () will solve the following QP.

$$\begin{aligned} u_{\text{opt}} &= \underset{u \in \mathbb{R}^2}{\operatorname{argmin}} \|u - u_{\text{ref}}\|^2 \\ \text{subject to } &A_{\text{barrier}}u + b_{\text{barrier}} \leq 0 \end{aligned} \quad (6)$$

Question report 3.2 (0.5 points). Comparing (6) with (5), and knowing that there are 4 obstacles in the environment, give the dimensions of A_{barrier} and b_{barrier} .

Question provided 3.2. Write a class providing a way to evaluate the CLF and to compute u^* .

File name: `me570_potential.py`

Class name: `Clfcbf_Control`

Description: A class implementing a CLF-CBF-based control framework.

method name: `__init__`

Description: Save the arguments to internal attributes, and create an attribute `attractive` with an object of class `Attractive` using the argument `potential`.

Input arguments

- `world` (type `SphereWorld`), `potential` (type `dict`): same as in the class `Total` (`_`).

method name: `function`

Description: Evaluate the CLF (i.e., `self.attractive.eval()`) at the given input.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): point at which to compute u^* .

Question **code** 3.1. Write a method to compute u^* .

Class name: `Clfcbf_Control`

File name: `me570_potential.py`

Method name: `control`

Description: Compute u^* according to (5).

Input arguments

- `x_eval` (dim. $[2 \times 1]$): point at which to compute u^* .

Output arguments

- `u_opt` (dim. $[2 \times 1]$): The solution to the QP (5).

Requirements: This function should use `qp_supervisor` (`_`) from Question provided 3.1, by building A_{barrier} , b_{barrier} according to ♠ in (5b), using

`ch = self.potential['repulsive_weight']` .

Question **provided** 3.3.

File name: `me570_hw3.py`

Method name: `clfcbf_control_test_singlesphere`

Description: Use the provided function `Grid.plot_threshold` (`_`) to visualize the CLF-CBF control field around a single filled-in sphere with the goal set at $x = \begin{bmatrix} 0 \\ -6 \end{bmatrix}$.

Question **report** 3.3 (0.5 points). Run the function from provided 3.3 above, and include the produced image in the report. You should see that the field induces a flow that goes around and has a stable equilibrium around the point $\begin{bmatrix} 0 \\ -6 \end{bmatrix}$.

Question **report** 3.4 (0.5 points). Explain why the field is essentially zero at the origin (i.e., at the top of the edge of the sphere). To get full points, your explanation should be based on 5 and what we have seen in class.

Question report 3.5. This question is analogous to report 2.3, but using the CLF-CBF framework instead of the gradient of the potential.

File name: `me570_hw3.py`

Method name: `clfcbf_run_plot_test`

Description: Use the function `Planner.run_plot(.)` to run the planner based on the CLF-CBF framework, and show the results for one combination of `repulsive_weight` and `epsilon` that makes the planner work reliably.

Requirements: The function should follow these steps *for each goal*:

- 1) Load the sphere world in the variable `world`
- 2) Create the dict `potential`.
- 3) Create an object of class `Clfcbf_Control`
- 4) Create an object of class `Planner`
- 5) Call `Planner.run_plot(.)`

Use the following for the arguments in the various steps above:

- `function`: use `Clfcbf_Control.function`;
- `control`: use `Clfcbf_Control.control`;
- `nb_steps`: use `20` (but increase if necessary);
- `potential['shape']`: use `'conic'`.

Include all figures in your report, trying to have all of them in one or two pages, for ease of comparison.

Question report 3.6. Use the function `field_plot_threshold(.)` to visualize the control field $u^*(x)$ for the value of `['repulsiveWeight']` included in the previous question (report 2.3). Make sure to superimpose the field on top of the `SphereWorld` map. Note that the computation of $u^*(x)$ will take significantly longer than just the gradient, hence you might want to reduce the size of the grid passed to `field_plot_threshold(.)` (e.g., use a 10×10 grid). Include the images in your report.

Question report 3.7. Comment on the trade-off between traditional gradient-based methods and a CLF-CBF formulation.

Question optional 3.1. Repeat the previous questions with `['shape']` set to `'quadratic'`.

Problem 4: Jacobian-based Inverse Kinematics (IK) for the two-link manipulator

In this problem, you will combine the attractive potential from Problem 2 with the Jacobian of the two-link manipulator from the previous homework to create a way to move the end effector of the manipulator to a specific configuration by acting on its joint angles. This procedure is commonly called Inverse Kinematics.

Question report 4.1. Recall from Homework 2 that we denote the position of the end

effector in the world reference frame as ${}^{\mathcal{W}}p_{\text{eff}}$. Based on the results of Homework 2, write an expression for the Jacobian matrix J such that $\frac{d}{dt}({}^{\mathcal{W}}p_{\text{eff}}) = J(\theta)\dot{\theta}$ (this was optional in Homework 2, but it is now mandatory). To obtain full marks, derive the expression in matrix form using the derivative of rotations (which involves rotations and skew-symmetric matrices), and show all the steps of the derivation.

Question **code** **4.1.** Implement a function for computing the matrix J .

Class name: `TwoLink`

File name: `me570_robot.py`

Description: See description from previous homework assignments.

Method name: `jacobian_matrix`

Description: Compute the matrix representation of the Jacobian of the position of the end effector with respect to the joint angles as derived in Question report 4.1.

Input arguments

- `theta` (dim. $[2 \times 1]$) An array containing $\begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$, the two joint angles for the two-link manipulator.

Output arguments

- `jtheta` (dim. $[2 \times 2]$): The matrix $J(\theta)$ defined in Question report 4.1.

If you already solved this question in Homework 2, please copy the relevant portion of your previous report here.

Question **optional** **4.1.** Compare the results of `Jtheta*thetaDot`, where `[Jtheta]=twolink_jacobianMatrix(theta)`, with the results of `Twolink.jacobian(theta,thetaDot)`, for arbitrary values of `theta` and `thetaDot` (they should be the same).

Question **provided** **4.1.** In this question you will adapt the potential planner from Problem 2 to work with the two-link manipulator, by pulling back the total potential and its gradient from \mathbb{R}^2 to the configuration space of the two-link manipulator.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Description: Combines attractive and repulsive potentials

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `world` (type `SphereWorld`): an instance of the class `SphereWorld` (see the corresponding description for details).
- `potential` (dim. $[1 \times 1]$, type `struct`): the dictionary with fields `'x_goal'`, `'shape'` and `'repulsive_weight'` previously described. This function uses the field `'repulsive_weight'`, and then the structure is passed to an instance of `Attractive`.

Question code 4.2. Now it is time to implement the evaluation of the total potential and of its gradient for the `TwoLinkPotential` class.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `eval`

Description: Compute the potential U pulled back through the kinematic map of the two-link manipulator, i.e., $U({}^{\mathcal{W}}p_{\text{eff}}(\vec{\theta}))$, where U is defined as in Question code 2.3, and ${}^{\mathcal{W}}p_{\text{eff}}(\theta)$ is the position of the end effector in the world frame as a function of the joint angles $\vec{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$.

Input arguments

- `theta_eval` (dim. $[2 \times 1]$): the joint angles θ at which to evaluate the potential.

Output arguments

- `u_eval_theta`: The value of the potential U evaluated at `x_eval`.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `grad`

Description: Compute the gradient of the potential U pulled back through the kinematic map of the two-link manipulator, i.e., $\nabla_{\vec{\theta}} U({}^{\mathcal{W}}p_{\text{eff}}(\vec{\theta}))$.

Input arguments

- `theta_eval` (dim. $[2 \times 1]$): an array containing the two joint angles $\begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$ at which the gradient should be evaluated.

Output arguments

- `grad_u_eval_theta` (dim. $[2 \times 1]$): The gradient of the pulled-back potential, evaluated at `theta_eval`.

Note that for computing the gradient of the pulled-back function, you will need to use the Jacobian matrix given by `Twolink.jacobianMatrix(_)` (see also the material from class).

Question provided 4.2. Make a function `TwoLink.plot_animate(theta_path)` that uses the function `Twolink.plot(_)` from Homework 2 to show the sequence of configurations on the path specified by `theta_path`. Make all the drawing overlap (so that you can see all the configurations in a single picture). If you feel there are too many configurations, you can draw only a subset of the configurations (e.g., one configuration every 15).

Question provided 4.3. Implement a function that tests the potential planner applied to the two-link manipulator. This is similar to Question ??, except that we are planning the trajectory of the end effector of the manipulator, instead of a free point.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `run_plot`

Description: This function performs the same steps as `Planner.run_plot(_)` in Question ??, except for the following:

- In step 2)(b)iv: `U` should be set to `TwoLinkPotential.eval`, and `control` to the negative of `TwoLinkPotential.grad`.
- In step 2)(b)iv: Use the contents of the variable `thetaStart` instead of `xStart` to initialize the planner, and use only the second goal `x_goal[:,1]`.
- In step 2)(b)v: Use `Twolink.plot_animate(_)` to plot a decimated version of the results of the planner. Note that the output `xPath` from `Potential.planner(_)` will really contain a sequence of joint angles, rather than a sequence of 2-D points. Plot only every 5th or 10th column of `xPath` (e.g., use `xPath(:,1:5:end)`). To avoid clutter, plot a different figure for each start.

Input arguments

- `epsilon`, `nb_steps`: same as in `Planner(_)`.

Question report 4.2. Show the results of `twolink_planner_runPlot(_)` for (one) combination of `['repulsive_weight']` and `epsilon` that makes the planner work reliably. For the argument `nb_steps`, use `400`, and for `potential['shape']`, use `'quadratic'`. Note that `U`, and `control` are set by the function, so they do not need to be set in the argument. In your report, try to have all figures on the same one or two pages for ease of comparison.

If it is too difficult to get the planner work reliably for all goal configurations, it is fine to include the results using different parameters for different goals, or just for a few out of the five goals.

Note that, in this problem, we are considering only collisions between the spheres and the end effector of the manipulator; we are not considering collisions between the spheres

and the links of the manipulator; in other words, it is normal to have overlap between the manipulator and the spheres, as long as the end effector is not inside an obstacle.

Hint for question report 2.3: For the argument `control`, you can use `@(x) -potential_totalGrad(x,world,potential)`. If everything works, the plot of `URep` should be monotonically decreasing, at least for some very small `epsilon`, even if the planner does not get to the goal. As a rule of thumb, the values of `['repulsiveWeight']` and `epsilon` can be significantly different between the `'conic'` and `'quadratic'` shapes, and when `['repulsiveWeight']` is increased, the value of `epsilon` that works tends to decrease (given our discussion in class, can you explain this?) Generally, if the planner “goes crazy”, it is because either the repulsive gradient or the step size are too large. If the planner never enters the region of influence of the obstacles, the weight for the repulsive term is too large. If the planner does not make enough progress, the step size is too small.

Hint for question report 4.1: Note that expressions such as $av + bw$, where $a, b \in \mathbb{R}$ are scalars, and $v, w \in \mathbb{R}^2$ are vectors, can be equivalently written as the matrix-vector multiplication $\begin{bmatrix} v & w \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix}$.

Hint for question code 2.3: You can refer to `as` as an example of how to iterate over the `SphereWorld` object.

Hint for question report 3.2: Knowing that $u \in \mathbb{R}^{2 \times 1}$, the dimensions of A_{barrier} and b_{barrier} should be consistent with matrix-vector multiplication rules.